

# TP-ML-NANA-ROMARIC-v1

September 16, 2021

## 1 TRAVAUX PRATIQUES

### 1.0.1 COURS DE MACHINE LEARNING - UNIVERSITE VIRTUELLE DU BURKINA FASO - MASTER FD & IA

- ETUDIANT : NANA SIDWENDLUAN ROMARIC
- ENSEIGNANT : MADAME BIRBA ELIANE

#### 1.0.2 1. Chargement des données

```
[1]: #importing pyspark
import pyspark

#importing sparksession
from pyspark.sql import SparkSession

from pyspark.sql.functions import *

from pyspark.ml.classification import LogisticRegression
from pyspark.ml.evaluation import BinaryClassificationEvaluator

from pyspark.ml.classification import DecisionTreeClassifier

from pyspark.ml.classification import RandomForestClassifier
```

```
[2]: #creating a sparksession object and providing appName
spark=SparkSession.builder.master("local").appName("tp").getOrCreate()
```

```
[3]: datadft = spark.read.format("csv").options(header=True,inferSchema=True).
    ↳load("data/ccdefault.csv")
```

#### 1.0.3 2. Analyse exploratoire

```
[4]: datadft.printSchema()
```

```
root
 |-- ID: integer (nullable = true)
 |-- LIMIT_BAL: integer (nullable = true)
```

```

|-- SEX: integer (nullable = true)
|-- EDUCATION: integer (nullable = true)
|-- MARRIAGE: integer (nullable = true)
|-- AGE: integer (nullable = true)
|-- PAY_0: integer (nullable = true)
|-- PAY_2: integer (nullable = true)
|-- PAY_3: integer (nullable = true)
|-- PAY_4: integer (nullable = true)
|-- PAY_5: integer (nullable = true)
|-- PAY_6: integer (nullable = true)
|-- BILL_AMT1: integer (nullable = true)
|-- BILL_AMT2: integer (nullable = true)
|-- BILL_AMT3: integer (nullable = true)
|-- BILL_AMT4: integer (nullable = true)
|-- BILL_AMT5: integer (nullable = true)
|-- BILL_AMT6: integer (nullable = true)
|-- PAY_AMT1: integer (nullable = true)
|-- PAY_AMT2: integer (nullable = true)
|-- PAY_AMT3: integer (nullable = true)
|-- PAY_AMT4: integer (nullable = true)
|-- PAY_AMT5: integer (nullable = true)
|-- PAY_AMT6: integer (nullable = true)
|-- DEFAULT: integer (nullable = true)

```

```
[5]: datadft.count()
```

```
[5]: 30000
```

```
[6]: datadft.show(5)
```

```

+---+-----+---+-----+-----+---+-----+-----+-----+-----+-----+---+
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+---+-----+-----+-----+-----+
| ID|LIMIT_BAL|SEX|EDUCATION|MARRIAGE|AGE|PAY_0|PAY_2|PAY_3|PAY_4|PAY_5|PAY_6|BI
LL_AMT1|BILL_AMT2|BILL_AMT3|BILL_AMT4|BILL_AMT5|BILL_AMT6|PAY_AMT1|PAY_AMT2|PAY_
AMT3|PAY_AMT4|PAY_AMT5|PAY_AMT6|DEFAULT|
+---+-----+---+-----+-----+---+-----+-----+-----+-----+-----+---+
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+---+-----+-----+-----+-----+
| 1| 20000| 2| 2| 1| 24| 2| 2| -1| -1| -2| -2|
3913| 3102| 689| 0| 0| 0|
0| 0| 0| 0| 1|
| 2| 120000| 2| 2| 2| 26| -1| 2| 0| 0| 0| 2|
2682| 1725| 2682| 3272| 3455| 3261| 0| 1000|
1000| 1000| 0| 2000| 1|
| 3| 90000| 2| 2| 2| 34| 0| 0| 0| 0| 0| 0|
29239| 14027| 13559| 14331| 14948| 15549| 1518| 1500|

```

```

1000|    1000|    1000|    5000|    0|
| 4|    50000| 2|    2|    1| 37|    0|    0|    0|    0|    0|
46990|    48233|    49291|    28314|    28959|    29547|    2000|    2019|
1200|    1100|    1069|    1000|    0|
| 5|    50000| 1|    2|    1| 57|   -1|    0|   -1|    0|    0|
8617|    5670|    35835|    20940|    19146|    19131|    2000|    36681|
10000|    9000|    689|    679|    0|
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
----+-----+-----+-----+-----+
only showing top 5 rows

```

```
[7]: datadft.describe(datadft.columns).show()
```

```

+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+
+-----+-----+-----+-----+
-----+-----+-----+-----+
-----+-----+-----+-----+
+-----+-----+-----+-----+
|summary|          ID|      LIMIT_BAL|          SEX|
EDUCATION|      MARRIAGE|          AGE|      PAY_0|
PAY_2|      PAY_3|      PAY_4|      PAY_5|
PAY_6|      BILL_AMT1|      BILL_AMT2|      BILL_AMT3|      BILL_AMT4|
BILL_AMT5|      BILL_AMT6|      PAY_AMT1|      PAY_AMT2|
PAY_AMT3|      PAY_AMT4|      PAY_AMT5|      PAY_AMT6|
DEFAULT|
+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+
+-----+-----+-----+-----+
-----+-----+-----+-----+
-----+-----+-----+-----+
+-----+-----+-----+-----+
| count|      30000|      30000|      30000|
30000|      30000|      30000|      30000|
30000|      30000|      30000|      30000|
30000|      30000|      30000|      30000|      30000|
30000|      30000|      30000|      30000|      30000|
30000|      30000|      30000|      30000|
| mean|      15000.5|167484.32266666667|1.6037333333333332|1.8531333333333
333|1.5518666666666667|      35.4855|
-0.0167|-0.13376666666666667|      -0.1662|-0.22066666666666668|
-0.2662|      -0.2911|      51223.3309|49179.07516666667|      47013.1548|
43262.948966666666|40311.40096666667|      38871.7604|      5663.5805|
5921.1635|      5225.6815| 4826.076866666666|
4799.3876333333334|5215.502566666667|      0.2212|
| stddev|8660.398374208891|129747.66156720246|0.4891291960902602|0.7903486597207

```

```

269|0.5219696006132467|9.217904068090155|1.1238015279973335|
1.1971859730345495|1.1968675684465686| 1.1691386224023357|1.1331874060027525|1.
149987625607897|73635.86057552966|71173.76878252832|69349.38742703677|64332.8561
33916444|60797.15577026471|59554.1075367459|16563.28035402577|23040.870402057186
|17606.96146980311|15666.159744032062|15278.305679144742|17777.46577543531|0.415
06180569093254|

```

```

|    min|                1|                10000|                1|
0|                0|                21|                -2|                -2|
-2|                -2|                -2|                -2|                -165580|
-69777|                -157264|                -170000|                -81334|                -339603|
0|                0|                0|                0|                0|
0|                0|
|    max|                30000|                1000000|                2|
6|                3|                79|                8|                8|
8|                8|                8|                8|                964511|
983931|                1664089|                891586|                927171|                961664|
873552|                1684259|                896040|                621000|
426529|                528666|                1|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+

```

```

[8]: datadft.describe("LIMIT_BAL", "BILL_AMT1", "PAY_AMT1", "BILL_AMT2", "PAY_AMT2").
      ↪ show()

```

```

+-----+-----+-----+-----+-----+
+-----+-----+
|summary|    LIMIT_BAL|    BILL_AMT1|    PAY_AMT1|
BILL_AMT2|    PAY_AMT2|
+-----+-----+-----+-----+-----+
+-----+-----+
|  count|    30000|    30000|    30000|
30000|    30000|
|   mean|167484.3226666667|    51223.3309|
5663.5805|49179.0751666667|    5921.1635|
|  stddev|129747.66156720246|73635.86057552966|16563.28035402577|71173.7687825283
2|23040.870402057186|
|   min|    10000|   -165580|         0|
-69777|         0|
|   max|    1000000|    964511|    873552|
983931|    1684259|
+-----+-----+-----+-----+-----+
+-----+-----+

```

```
[9]: datadft.describe("LIMIT_BAL", "AGE").show()
```

```
+-----+-----+-----+
|summary|          LIMIT_BAL|          AGE|
+-----+-----+-----+
|  count|          30000|          30000|
|   mean|167484.32266666667|          35.4855|
| stddev|129747.66156720246|9.217904068090155|
|   min|          10000|          21|
|   max|         1000000|          79|
+-----+-----+-----+
```

```
[10]: datadft.groupBy("SEX").count().orderBy(asc("count")).show() # SEXE ? HOMME=1
      ↪ FEMME=2
```

```
+---+-----+
|SEX|count|
+---+-----+
|  1|11888|
|  2|18112|
+---+-----+
```

```
[11]: datadft.groupBy("DEFAULT").count().orderBy(asc("count")).show() # DÉFAUT DE
      ↪ PAIEMENT ? OUI=1 NON=0
```

```
+-----+-----+
|DEFAULT|count|
+-----+-----+
|      1| 6636|
|      0|23364|
+-----+-----+
```

```
[12]: datadft.groupBy(['DEFAULT', 'SEX']).count().orderBy(asc("DEFAULT")).show()
      # repartition de la variable cible en fonction du sexe
```

```
+-----+---+-----+
|DEFAULT|SEX|count|
+-----+---+-----+
|      0|  1| 9015|
|      0|  2|14349|
|      1|  2| 3763|
|      1|  1| 2873|
+-----+---+-----+
```

```
[13]: datadft.groupBy(['DEFAULT', 'EDUCATION']).count().orderBy(asc("DEFAULT")).show()
# repartition de la variable cible en fonction du niveau d'instruction
```

DEFAULT	EDUCATION	count
0	0	14
0	1	8549
0	5	262
0	6	43
0	2	10700
0	3	3680
0	4	116
1	2	3330
1	3	1237
1	1	2036
1	5	18
1	4	7
1	6	8

Les clients qui ont un niveau d'instruction plus élevé sont les plus nombreux à avoir un défaut de paiement. D'après l'énoncé, la variable EDUCATION a pour valeur { 1 = graduate school; 2 = university; 3 = high school; 4 = others }. Nous constatons cependant qu'il y a des lignes où la variable EDUCATION a des valeurs plus grandes que 4. Dans ce fichier nous conservons ces lignes, puis nous appliquerons un traitement plus tard dans un autre fichier.

```
[14]: datadft.groupBy(['DEFAULT', 'MARRIAGE']).count().orderBy(asc("DEFAULT")).show()
# repartition de la variable cible en fonction du statut matrimonial
```

DEFAULT	MARRIAGE	count
0	0	49
0	1	10453
0	2	12623
0	3	239
1	0	5
1	2	3341
1	1	3206
1	3	84

D'après l'énoncé, la variable MARRIAGE a pour valeur { 1 = married; 2 = single; 3 = others }. Nous constatons cependant qu'il y a des lignes où la variable MARRIAGE a une valeur 0 dans certaines lignes. Dans ce fichier nous conservons ces lignes, puis nous appliquerons un traitement plus tard dans un autre fichier.

```
[15]: datadft.groupBy(['DEFAULT', 'AGE']).count().orderBy(asc("DEFAULT")).show()
```

```
+-----+---+-----+
|DEFAULT|AGE|count|
+-----+---+-----+
|      0| 42|  609|
|      0| 27| 1164|
|      0| 39|  755|
|      0| 58|   91|
|      0| 71|    3|
|      0| 28| 1123|
|      0| 56|  129|
|      0| 50|  310|
|      0| 22|  391|
|      0| 31|  988|
|      0| 67|   11|
|      0| 40|  683|
|      0| 57|   95|
|      0| 32|  933|
|      0| 60|   44|
|      0| 73|    1|
|      0| 65|   19|
|      0| 70|    8|
|      0| 48|  362|
|      0| 25|  884|
+-----+---+-----+
```

only showing top 20 rows

```
[16]: datadft.groupBy(['DEFAULT', 'LIMIT_BAL']).count().orderBy(desc("DEFAULT")).
      ↪orderBy(desc("count")).show()
```

```
+-----+-----+-----+
|DEFAULT|LIMIT_BAL|count|
+-----+-----+-----+
|      0|   50000| 2480|
|      0|   20000| 1278|
|      0|  200000| 1258|
|      0|   80000| 1204|
|      0|   30000| 1042|
|      0|  150000|  923|
|      1|   50000|  885|
|      0|  180000|  819|
|      0|  100000|  776|
|      0|  360000|  727|
|      1|   20000|  698|
|      0|  500000|  641|
|      0|  230000|  624|
```

```

|      0|    210000|    613|
|      0|     60000|    592|
|      0|   140000|    579|
|      0|   130000|    572|
|      1|    30000|    568|
|      0|   160000|    557|
|      0|   120000|    547|
+-----+-----+-----+
only showing top 20 rows

```

```
[17]: datadft.createOrReplaceTempView("dataView")
      spark.sql("SELECT DEFAULT, avg(LIMIT_BAL) AS BALANCE FROM dataView GROUP BY_
      ↪DEFAULT ORDER BY BALANCE DESC").show()
```

```

+-----+-----+-----+
|DEFAULT|          BALANCE|
+-----+-----+-----+
|      0|178099.72607430234|
|      1|130109.65641952984|
+-----+-----+-----+

```

```
[18]: datadft
```

```
[18]: DataFrame[ID: int, LIMIT_BAL: int, SEX: int, EDUCATION: int, MARRIAGE: int, AGE:
int, PAY_0: int, PAY_2: int, PAY_3: int, PAY_4: int, PAY_5: int, PAY_6: int,
BILL_AMT1: int, BILL_AMT2: int, BILL_AMT3: int, BILL_AMT4: int, BILL_AMT5: int,
BILL_AMT6: int, PAY_AMT1: int, PAY_AMT2: int, PAY_AMT3: int, PAY_AMT4: int,
PAY_AMT5: int, PAY_AMT6: int, DEFAULT: int]
```

### 1.0.4 3. Preparation des données

Renommons la colonne DEFAULT en label

```
[19]: # datadft_bis = datadft.
      ↪withColumn("ID", "LIMIT_BAL", "SEX", "EDUCATION", "MARRIAGE", "AGE", "PAY_0", "PAY_2", "PAY_3", "PAY_
renamedDatadft = datadft.withColumnRenamed("DEFAULT", "label")
```

```
[20]: renamedDatadft.show(5)
```

```

+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| ID|LIMIT_BAL|SEX|EDUCATION|MARRIAGE|AGE|PAY_0|PAY_2|PAY_3|PAY_4|PAY_5|PAY_6|BI
LL_AMT1|BILL_AMT2|BILL_AMT3|BILL_AMT4|BILL_AMT5|BILL_AMT6|PAY_AMT1|PAY_AMT2|PAY_
AMT3|PAY_AMT4|PAY_AMT5|PAY_AMT6|label|
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```



```

-----+-----+-----+-----+-----+
| 1|    20000| 2|        2|        1| 24|    2|    2|   -1|   -1|   -2|   -2|
3913|    3102|    689|        0|        0|    0|    0|    0|    689|
0|    0|    0|    0|    1|
| 2|   120000| 2|        2|        2| 26|   -1|    2|    0|    0|    0|    2|
2682|    1725|   2682|   3272|   3455|   3261|    0|   1000|
1000|    1000|    0|   2000|    1|
| 3|    90000| 2|        2|        2| 34|    0|    0|    0|    0|    0|    0|
29239|   14027|   13559|   14331|   14948|   15549|   1518|   1500|
1000|    1000|   1000|   5000|    0|
| 4|    50000| 2|        2|        1| 37|    0|    0|    0|    0|    0|    0|
46990|   48233|   49291|   28314|   28959|   29547|   2000|   2019|
1200|    1100|   1069|   1000|    0|
| 5|    50000| 1|        2|        1| 57|   -1|    0|   -1|    0|    0|    0|
8617|    5670|   35835|   20940|   19146|   19131|   2000|   36681|
10000|    9000|    689|    679|    0|
-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+

```

only showing top 5 rows

```

[21]: # colonne des etiquettes
      collLabel = "label"

      # colonne numerique
      colNum = [col for col in renamedDatadft.columns if col!= collLabel]

```

```

[22]: colNum

```

```

[22]: ['ID',
      'LIMIT_BAL',
      'SEX',
      'EDUCATION',
      'MARRIAGE',
      'AGE',
      'PAY_0',
      'PAY_2',
      'PAY_3',
      'PAY_4',
      'PAY_5',
      'PAY_6',
      'BILL_AMT1',
      'BILL_AMT2',
      'BILL_AMT3',
      'BILL_AMT4',
      'BILL_AMT5',

```

```
'BILL_AMT6',
'PAY_AMT1',
'PAY_AMT2',
'PAY_AMT3',
'PAY_AMT4',
'PAY_AMT5',
'PAY_AMT6']
```

```
[23]: from pyspark.ml.feature import VectorAssembler, StandardScaler

va = VectorAssembler().setInputCols(colNum).
    ↪setOutputCol("to_be_scaled_features")

featuredDatadft = va.transform(renamedDatadft)

scaler = StandardScaler().setInputCol("to_be_scaled_features").
    ↪setOutputCol("features")

dataset = scaler.fit(featuredDatadft).transform(featuredDatadft).
    ↪select("features", "label")

dataset.show(5)
```

```
+-----+-----+
|          features|label|
+-----+-----+
|[1.15468129385139...|    1|
|[2.30936258770278...|    1|
|[3.46404388155417...|    0|
|[4.61872517540556...|    0|
|[5.77340646925695...|    0|
+-----+-----+
only showing top 5 rows
```

#### 1.0.5 4. Application des modèles

```
[24]: trainSet, testSet = dataset.randomSplit([0.8,0.2])
```

```
[25]: trainSet.count()
```

```
[25]: 24010
```

```
[26]: testSet.count()
```

```
[26]: 5990
```

```
[27]: trainSet.show(5)
```

```
+-----+-----+
|          features|label|
+-----+-----+
|(24,[0,1,2,3,4,5,...|    0|
|(24,[0,1,2,3,4,5,...|    0|
|(24,[0,1,2,3,4,5,...|    1|
|(24,[0,1,2,3,4,5,...|    1|
|(24,[0,1,2,3,4,5,...|    1|
+-----+-----+
only showing top 5 rows
```

## 1.0.6 4.1 Logistic Regression

```
[28]: from pyspark.ml.classification import LogisticRegression
      from pyspark.ml.evaluation import BinaryClassificationEvaluator

      lr = LogisticRegression(maxIter=100, regParam=0.0001, elasticNetParam=0.1)
      lrModel = lr.fit(trainSet)
```

```
[29]: # trainingSummary = lrModel.summary
      print("Coefficients: " + str(lrModel.coefficients))
      print("Intercept: " + str(lrModel.intercept))
```

```
Coefficients: [-0.021741799077511652,-0.10177122033887191,-0.04521580458898812,-
0.07114937293373588,-0.07936493858026904,0.06967836813214995,0.6251280383863709,
0.11493582924386118,0.06944719240978961,0.05505572695441151,0.037448435283436596
,0.0005579010558137603,-0.3227225426356211,-0.03194435843587904,0.11429461888777
985,0.05079451862353235,0.01823735451185488,0.05562332444776182,-0.1562284789069
5163,-0.2291329735750258,-0.08667448207609935,-0.05519722837168079,-0.0664881969
8795533,-0.045074589359989314]
Intercept: -0.6853446776524614
```

```
[30]: summary = lrModel.summary
      print("training set AreaUnderROC: ",summary.areaUnderROC)
      summary.roc.show()
      summary.pr.show()
```

```
training set AreaUnderROC: 0.7221510619656176
+-----+-----+
|          FPR|          TPR|
+-----+-----+
|          0.0|          0.0|
|4.812834224598930...|0.002824858757062147|
|0.001069518716577...|0.005273069679849341|
|0.001497326203208556|0.008286252354048965|
```

```
|0.001871657754010...|0.011487758945386064|
|0.002085561497326...|0.015254237288135594|
|0.002139037433155...|0.019585687382297552|
|0.002459893048128...|0.022975517890772127|
|0.002513368983957...|0.027306967984934087|
|0.002727272727272...|0.031073446327683617|
|0.003262032085561...| 0.03370998116760829|
|0.003475935828877...|0.037476459510357815|
|0.003850267379679...| 0.04067796610169491|
|0.004278074866310...|0.043691148775894535|
|0.004652406417112299| 0.04689265536723164|
|0.005026737967914439|0.050094161958568736|
|0.005240641711229946| 0.05386064030131827|
|0.005508021390374332| 0.05743879472693032|
|0.005989304812834224|0.060263653483992465|
|0.006310160427807487| 0.06365348399246705|
+-----+
```

only showing top 20 rows

```
+-----+
|          recall|          precision|
+-----+
|          0.0|          0.625|
|0.002824858757062147|          0.625|
|0.005273069679849341|0.5833333333333334|
|0.008286252354048965|0.6111111111111112|
|0.011487758945386064|0.6354166666666666|
|0.015254237288135594|          0.675|
|0.019585687382297552|0.7222222222222222|
|0.022975517890772127|0.7261904761904762|
|0.027306967984934087|0.7552083333333334|
|0.031073446327683617|0.7638888888888888|
| 0.03370998116760829|0.7458333333333333|
|0.037476459510357815|0.7537878787878788|
| 0.04067796610169491|          0.75|
|0.043691148775894535|0.7435897435897436|
| 0.04689265536723164|0.7410714285714286|
|0.050094161958568736|0.7388888888888889|
| 0.05386064030131827|0.7447916666666666|
| 0.05743879472693032|0.7475490196078431|
|0.060263653483992465|0.7407407407407407|
| 0.06365348399246705|0.7412280701754386|
+-----+
```

only showing top 20 rows

```
[31]: lr_predictions = lrModel.transform(testSet)
lr_predictions.select("prediction", "label", "features").show(25)

lr_evaluator = BinaryClassificationEvaluator()
print('Test Area Under ROC', lr_evaluator.evaluate(lr_predictions))
```

```
+-----+-----+-----+
|prediction|label|          features|
+-----+-----+-----+
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
+-----+-----+-----+
```

only showing top 25 rows

Test Area Under ROC 0.7324437493693866

```
[32]: lr_evaluator2 = BinaryClassificationEvaluator()
accuracy = lr_evaluator2.evaluate(lr_predictions)
print("Accuracy = %s" % (accuracy))
print("Test Error = %s" % (1.0 - accuracy))
```

Accuracy = 0.7324437493693866  
Test Error = 0.2675562506306134

### 1.0.7 4.2 Decision Tree

```
[34]: from pyspark.ml.classification import DecisionTreeClassifier
      from pyspark.ml.evaluation import BinaryClassificationEvaluator

      dt2 = DecisionTreeClassifier().setLabelCol("label").setFeaturesCol("features").
      ↪ setMaxDepth(8)

      dtModel2 = dt2.fit(trainSet)

      # make predictions on the test data
      dt_predictions2 = dtModel2.transform(testSet)
      dt_predictions2.select("prediction", "label", "features").show(25)

      # evaluate the model
      dt_evaluator2 = BinaryClassificationEvaluator()
      print("Test Area Under ROC: " + str(dt_evaluator2.evaluate(dt_predictions2,
      ↪ {dt_evaluator2.metricName: "areaUnderROC"})))
```

```
+-----+-----+-----+
|prediction|label|          features|
+-----+-----+-----+
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
+-----+-----+-----+
```

only showing top 25 rows

Test Area Under ROC: 0.37598343203226675

### 1.0.8 4.3 Random Forest

```
[35]: from pyspark.ml.classification import RandomForestClassifier
      from pyspark.ml.evaluation import BinaryClassificationEvaluator

      rfClassifier = RandomForestClassifier().setLabelCol("label").
        ↪ setFeaturesCol("features")

      trainedModel = rfClassifier.fit(trainSet)

      # make predictions on the test data
      rf_predictions = trainedModel.transform(testSet)
      rf_predictions.select("prediction", "label", "features").show(25)

      # evaluate random forest model
      rf_evaluator = BinaryClassificationEvaluator()
      print("Test Area Under ROC: " + str(rf_evaluator.evaluate(rf_predictions,
        ↪ {rf_evaluator.metricName: "areaUnderROC"})))
```

prediction	label	features
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]

```
|      0.0|    0|(24,[0,1,2,3,4,5,...|
|      0.0|    1|(24,[0,1,2,3,4,5,...|
|      0.0|    0|(24,[0,1,2,3,4,5,...|
+-----+-----+-----+
only showing top 25 rows
```

Test Area Under ROC: 0.7687806898059399

```
[36]: rf_evaluator2 = BinaryClassificationEvaluator(labelCol="label")
      rf_accuracy = rf_evaluator2.evaluate(rf_predictions)
      print("Accuracy = %s" % (rf_accuracy))
      print("Test Error = %s" % (1.0 - rf_accuracy))
```

Accuracy = 0.7687806898059399

Test Error = 0.23121931019406006