

TP-ML-NANA-ROMARIC-v2

September 16, 2021

1 TRAVAUX PRATIQUES

1.0.1 COURS DE MACHINE LEARNING - UNIVERSITE VIRTUELLE DU BURKINA FASO - MASTER FD & IA

- ETUDIANT : NANA SIDWENDLUIAN ROMARIC
- ENSEIGNANT : MADAME BIRBA ELIANE

1.0.2 1. Chargement des données

```
[1]: #importing pyspark
import pyspark

#importing sparksession
from pyspark.sql import SparkSession

from pyspark.sql.functions import *

from pyspark.ml.classification import LogisticRegression
from pyspark.ml.evaluation import BinaryClassificationEvaluator

from pyspark.ml.classification import DecisionTreeClassifier

from pyspark.ml.classification import RandomForestClassifier
```

```
[2]: #creating a sparksession object and providing appName
spark=SparkSession.builder.master("local").appName("tp").getOrCreate()
```

Constatant les resultats mitigés obtenus dans le fichier précédent, nous avons retiré certaines lignes du dataset qui contenaient des anomalies. Il s'agit notamment des lignes où: * la variable EDUCATION a des valeurs plus grandes que 4 * la variable MARRIAGE a des valeurs 0 (ZÉRO) * les variables PAY_0,PAY_1,PAY_2,...,PAY_6 ont des valeurs -2

```
[3]: datadft_bis = spark.read.format("csv").options(header=True,inferSchema=True).
    ↪load("data/ccdefault-bis.csv")
```

1.0.3 2. Analyse exploratoire

```
[4]: datadft_bis.printSchema()
```

```
root
|-- ID: integer (nullable = true)
|-- LIMIT_BAL: integer (nullable = true)
|-- SEX: integer (nullable = true)
|-- EDUCATION: integer (nullable = true)
|-- MARRIAGE: integer (nullable = true)
|-- AGE: integer (nullable = true)
|-- PAY_0: integer (nullable = true)
|-- PAY_2: integer (nullable = true)
|-- PAY_3: integer (nullable = true)
|-- PAY_4: integer (nullable = true)
|-- PAY_5: integer (nullable = true)
|-- PAY_6: integer (nullable = true)
|-- BILL_AMT1: integer (nullable = true)
|-- BILL_AMT2: integer (nullable = true)
|-- BILL_AMT3: integer (nullable = true)
|-- BILL_AMT4: integer (nullable = true)
|-- BILL_AMT5: integer (nullable = true)
|-- BILL_AMT6: integer (nullable = true)
|-- PAY_AMT1: integer (nullable = true)
|-- PAY_AMT2: integer (nullable = true)
|-- PAY_AMT3: integer (nullable = true)
|-- PAY_AMT4: integer (nullable = true)
|-- PAY_AMT5: integer (nullable = true)
|-- PAY_AMT6: integer (nullable = true)
|-- DEFAULT: integer (nullable = true)
```

```
[5]: datadft_bis.count()
```

```
[5]: 53104
```

```
[6]: datadft_bis.where("ID is null").count()
```

```
[6]: 29946
```

```
[7]: datadft_bis.where("ID is not null").count()
```

```
[7]: 23158
```

```
[8]: datadft_bis_clean=datadft_bis.where("ID is not null")
```

```
[9]: datadft_bis_clean.count()
```

[9]: 23158

```
[12]: datadft_bis_clean.describe(datadft_bis_clean.columns).show()
```

	ID	LIMIT_BAL	SEX	
EDUCATION	MARRIAGE	AGE	PAY_0	
PAY_2	PAY_3	PAY_4	PAY_5	
PAY_6	BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4
BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	
PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6	
DEFAULT				
count	23158	23158	23158	
23158	23158	23158	23158	
23158	23158	23158	23158	
23158	23158	23158	23158	23158
23158	23158	23158	23158	
23158	23158	23158	23158	
23158				
mean	14916.514509024959	156289.99395457294	1.5915450384316434	1.851412039036
1861	1.5644269798773642	35.24509888591415	0.180326453061577	0.18874686933241214
0.17445375248294326	0.1332584851887037	0.0918041281630538	0.08826323516711287	6
1282.07107694965	59663.90612315399	57343.27321012177	53264.461222903534	49786.01
8870368775	47915.041195267295	6075.015243112532	6121.004836341653	
5562.976941013904	5089.545815700838	5049.283875982382		
5334.334312116763	0.2315398566370153			
stddev	8614.422528409888	127579.85957421701	0.4915586840970848	0.700640636966
5048	0.5187412618034684	9.291428745301381	0.9848565503867994	
1.0345749644615636	1.0223398941110264	0.9888780901255524	0.941911434746279	
0.9498271923186254	77602.34303734612	75193.40105589147	72368.62997241326	
68063.73018666815	64341.13974498193	63112.234644456585	16902.04758187008	20302.
039638491944	18130.689034411655	15541.648029188553	14947.882729376512	17297.4285
48142714	0.4218255978798094			
min	2	10000	1	

count	23158	23158
mean	156289.99395457294	35.24509888591415
stddev	127579.85957421701	9.291428745301381
min	10000	21
max	1000000	79

```
[15]: datadft_bis_clean.groupBy("SEX").count().orderBy(asc("count")).show() # SEXE ?
      ↪ HOMME=1 FEMME=2
```

SEX	count
1	9459
2	13699

```
[16]: datadft_bis_clean.groupBy("DEFAULT").count().orderBy(asc("count")).show() #
      ↪ DÉFAUT DE PAIEMENT ? OUI=1 NON=0
```

DEFAULT	count
1	5362
0	17796

```
[17]: datadft_bis_clean.groupBy(['DEFAULT', 'SEX']).count().orderBy(asc("DEFAULT")).
      ↪ show()
      # repartition de la variable cible en fonction du sexe
```

DEFAULT	SEX	count
0	1	7081
0	2	10715
1	2	2984
1	1	2378

```
[18]: datadft_bis_clean.groupBy(['DEFAULT', 'EDUCATION']).count().
      ↪ orderBy(asc("DEFAULT")).show()
      # repartition de la variable cible en fonction du niveau d'instruction
```

DEFAULT	EDUCATION	count
0	0	8
0	1	6085
0	2	8672
0	3	2962
0	4	69
1	2	2847
1	1	1479
1	3	1033
1	4	3

```
[19]: datadft_bis_clean.groupBy(['DEFAULT', 'MARRIAGE']).count().
      ↪orderBy(asc("DEFAULT")).show()
      # repartition de la variable cible en fonction du statut matrimonial
```

DEFAULT	MARRIAGE	count
0	1	7792
0	2	9809
0	3	195
1	2	2724
1	1	2564
1	3	74

```
[20]: datadft_bis_clean.groupBy(['DEFAULT', 'AGE']).count().orderBy(asc("DEFAULT")).
      ↪show()
```

DEFAULT	AGE	count
0	42	452
0	27	918
0	39	566
0	31	708
0	71	3
0	28	868
0	56	104
0	50	224
0	22	344
0	58	62
0	67	11

	0	40	491
	0	57	76
	0	32	699
	0	60	31
	0	73	1
	0	65	17
	0	70	6
	0	48	275
	0	25	722

+-----+-----+

only showing top 20 rows

```
[21]: datadft_bis_clean.groupBy(['DEFAULT', 'LIMIT_BAL']).count().
      ↪orderBy(desc("DEFAULT")).orderBy(desc("count")).show()
```

	DEFAULT	LIMIT_BAL	count
	0	50000	2137
	0	20000	1070
	0	80000	946
	0	30000	867
	0	200000	811
	1	50000	790
	0	150000	626
	0	100000	620
	1	20000	593
	0	180000	567
	0	60000	537
	1	30000	522
	0	140000	488
	0	70000	480
	0	500000	458
	0	130000	448
	0	210000	431
	0	120000	430
	0	230000	429
	0	360000	410

+-----+-----+

only showing top 20 rows

```
[22]: datadft_bis_clean.createOrReplaceTempView("dataView")
      ↪spark.sql("SELECT DEFAULT, avg(LIMIT_BAL) AS BALANCE FROM dataView GROUP BY_
      ↪DEFAULT ORDER BY BALANCE DESC").show()
```

+-----+-----+

	DEFAULT	BALANCE
0	168451.11260957518	
1	115928.32525177173	

```
[23]: datadft_bis_clean
```

```
[23]: DataFrame[ID: int, LIMIT_BAL: int, SEX: int, EDUCATION: int, MARRIAGE: int, AGE:
int, PAY_0: int, PAY_2: int, PAY_3: int, PAY_4: int, PAY_5: int, PAY_6: int,
BILL_AMT1: int, BILL_AMT2: int, BILL_AMT3: int, BILL_AMT4: int, BILL_AMT5: int,
BILL_AMT6: int, PAY_AMT1: int, PAY_AMT2: int, PAY_AMT3: int, PAY_AMT4: int,
PAY_AMT5: int, PAY_AMT6: int, DEFAULT: int]
```

1.0.4 3. Preparation des données

Renommons la colonne DEFAULT en label

```
[24]: # datadft_bis = datadft.
      ↪withColumn("ID", "LIMIT_BAL", "SEX", "EDUCATION", "MARRIAGE", "AGE", "PAY_0", "PAY_2", "PAY_3", "PAY_4", "PAY_5", "PAY_6", "BILL_AMT1", "BILL_AMT2", "BILL_AMT3", "BILL_AMT4", "BILL_AMT5", "BILL_AMT6", "PAY_AMT1", "PAY_AMT2", "PAY_AMT3", "PAY_AMT4", "PAY_AMT5", "PAY_AMT6", "label")
renamedDatadft_bis_clean = datadft_bis_clean.
      ↪withColumnRenamed("DEFAULT", "label")
```

```
[25]: renamedDatadft_bis_clean.show(5)
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	PAY_6	BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6	label
2	120000	2		2	26	-1	2	0	0	0	2	2682	1725	2682	3272	3455	3261	0	1000						
3	90000	2		2	34	0	0	0	0	0	0	29239	14027	13559	14331	14948	15549	1518	1500						
4	50000	2		2	37	0	0	0	0	0	0	46990	48233	49291	28314	28959	29547	2000	2019						
5	50000	1		2	57	-1	0	-1	0	0	0	8617	5670	35835	20940	19146	19131	2000	36681						
6	50000	1		2	37	0	0	0	0	0	0	64400	57069	57608	19394	19619	20024	2500	1815						


```

657|      1000|      1000|      800|      0|
+---+-----+---+-----+---+-----+---+-----+---+-----+---+-----+---+-----+---+
-----+-----+---+-----+---+-----+---+-----+---+-----+---+-----+---+-----+---+
---+-----+-----+-----+-----+-----+

```

only showing top 5 rows

```

[26]: # colonne des etiquettes
      collLabel = "label"

      # colonne numerique
      colNum = [col for col in renamedDatadft_bis_clean.columns if col!= collLabel]

```

```

[27]: colNum

```

```

[27]: ['ID',
      'LIMIT_BAL',
      'SEX',
      'EDUCATION',
      'MARRIAGE',
      'AGE',
      'PAY_0',
      'PAY_2',
      'PAY_3',
      'PAY_4',
      'PAY_5',
      'PAY_6',
      'BILL_AMT1',
      'BILL_AMT2',
      'BILL_AMT3',
      'BILL_AMT4',
      'BILL_AMT5',
      'BILL_AMT6',
      'PAY_AMT1',
      'PAY_AMT2',
      'PAY_AMT3',
      'PAY_AMT4',
      'PAY_AMT5',
      'PAY_AMT6']

```

```

[28]: from pyspark.ml.feature import VectorAssembler, StandardScaler

      va = VectorAssembler().setInputCols(colNum).
      ↪setOutputCol("to_be_scaled_features")

      featuredDatadft_bis_clean = va.transform(renamedDatadft_bis_clean)

```

```

scaler = StandardScaler().setInputCol("to_be_scaled_features").
↳setOutputCol("features")

dataset_bis = scaler.fit(featuredDatadft_bis_clean).
↳transform(featuredDatadft_bis_clean).select("features", "label")

dataset_bis.show(5)

```

```

+-----+-----+
|          features|label|
+-----+-----+
| [2.32168783618879...|    1|
| [3.48253175428319...|    0|
| [4.64337567237759...|    0|
| [5.80421959047199...|    0|
| [6.96506350856639...|    0|
+-----+-----+
only showing top 5 rows

```

1.0.5 4. Application des modèles

```
[29]: trainSetb, testSetb = dataset_bis.randomSplit([0.8,0.2])
```

```
[30]: trainSetb.count()
```

```
[30]: 18494
```

```
[31]: testSetb.count()
```

```
[31]: 4664
```

```
[32]: trainSetb.show(5)
```

```

+-----+-----+
|          features|label|
+-----+-----+
| (24, [0,1,2,3,4,5,...|    0|
| (24, [0,1,2,3,4,5,...|    0|
| (24, [0,1,2,3,4,5,...|    0|
| (24, [0,1,2,3,4,5,...|    0|
| (24, [0,1,2,3,4,5,...|    0|
+-----+-----+
only showing top 5 rows

```

1.0.6 4.1 Logistic Regression

```
[33]: from pyspark.ml.classification import LogisticRegression
      from pyspark.ml.evaluation import BinaryClassificationEvaluator

      lr_b = LogisticRegression(maxIter=100, regParam=0.0001, elasticNetParam=0.1)
      lrModel_b = lr_b.fit(trainSetb)
```

```
[34]: # trainingSummary = lrModel.summary
      print("Coefficients: " + str(lrModel_b.coefficients))
      print("Intercept: " + str(lrModel_b.intercept))
```

```
Coefficients: [-0.02094330588103058,-0.17998491808473413,-0.04570017072716318,-0
.03772318212736771,-0.08207283798231185,0.05980937514460258,0.6760420221760369,0
.11032980661501554,0.08180686709890893,0.04774788768784148,0.07102923985889595,0
.07978566469835958,-0.3711038032003112,0.2841025868205868,0.002833425120911844,0
.049224513407792644,0.015627669327713645,-0.015931441855878616,-0.15447222157692
014,-0.14494741912990822,-0.007457465858492866,-0.05608250525945112,-0.047916706
055974104,0.02873570956534439]
Intercept: -0.928546893254043
```

```
[35]: summary = lrModel_b.summary
      print("Training set areaUnderROC:",summary.areaUnderROC)
      summary.roc.show()
      summary.pr.show()
```

Training set areaUnderROC: 0.7524071908449781

```
+-----+-----+
|          FPR|          TPR|
+-----+-----+
|          0.0|          0.0|
|4.907459338194055E-4|0.002600472813238...|
|0.001051598429613...|0.004964539007092199|
|0.001542344363432...|0.007565011820330969|
|0.001752664049355...|0.011111111111111112|
|0.002243409983174425|0.013711583924349883|
|0.002383623107122...|0.017494089834515367|
|0.002734155916993...| 0.02056737588652482|
|0.003084688726864...| 0.02364066193853428|
|0.003295008412787437|0.027186761229314422|
|0.003715647784632...|0.030023640661938536|
|0.003996074032529445| 0.03333333333333333|
|0.004136287156477846|0.037115839243498816|
|0.004416713404374649| 0.04042553191489362|
|0.004837352776219...| 0.04326241134751773|
|0.005047672462142457| 0.04680851063829787|
| 0.00532809871003926|0.050118203309692674|
|0.005748738081884464|0.052955082742316785|
```

```
|0.005959057767807066|0.056501182033096925|
|0.006449803701626472| 0.0591016548463357|
+-----+
only showing top 20 rows
```

```
+-----+
|          recall|          precision|
+-----+
|          0.0|0.611111111111112|
|0.002600472813238...|0.611111111111112|
|0.004964539007092199|0.5833333333333334|
|0.007565011820330969|0.5925925925925926|
|0.01111111111111112|0.6527777777777778|
|0.013711583924349883|0.6444444444444445|
|0.017494089834515367|0.6851851851851852|
| 0.02056737588652482|0.6904761904761905|
| 0.02364066193853428|0.6944444444444444|
|0.027186761229314422|0.7098765432098766|
|0.030023640661938536|0.7055555555555556|
| 0.03333333333333333|0.7121212121212122|
|0.037115839243498816|0.7268518518518519|
| 0.04042553191489362|0.7307692307692307|
| 0.04326241134751773|0.7261904761904762|
| 0.04680851063829787|0.7333333333333333|
|0.050118203309692674|0.736111111111112|
|0.052955082742316785|0.7320261437908496|
|0.056501182033096925|0.7376543209876543|
| 0.0591016548463357|0.7309941520467836|
+-----+
only showing top 20 rows
```

```
[36]: # make predictions on the test data
lr_predictions_b = lrModel_b.transform(testSetb)
lr_predictions_b.select("prediction", "label", "features").show(25)

lr_evaluator_b = BinaryClassificationEvaluator()
print('Test Area Under ROC', lr_evaluator_b.evaluate(lr_predictions_b))
```

```
+-----+
|prediction|label|          features|
+-----+
|      0.0|    1|(24, [0,1,2,3,4,5,...|
|      0.0|    0|(24, [0,1,2,3,4,5,...|
|      0.0|    1|(24, [0,1,2,3,4,5,...|
|      0.0|    0|(24, [0,1,2,3,4,5,...|
|      0.0|    0|(24, [0,1,2,3,4,5,...|
```

0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	[2.32168783618879...
0.0	0	[0.00139301270171...
0.0	0	[0.00185735026895...
0.0	0	[0.00336644736247...
0.0	0	[0.00359861614609...
0.0	0	[0.00429512249694...
0.0	0	[0.00464337567237...
0.0	0	[0.00499162884780...
0.0	0	[0.00684897911675...
1.0	1	[0.00835807621027...

+-----+-----+-----+-----+
only showing top 25 rows

Test Area Under ROC 0.7364789966745253

1.0.7 4.2 Decision Tree

```
[40]: from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator

dt_b2 = DecisionTreeClassifier().setLabelCol("label").
    ↳setFeaturesCol("features").setMaxDepth(8)

dtModel_b2 = dt_b2.fit(trainSetb)

# make predictions on the test data
dt_predictions_b2 = dtModel_b2.transform(testSetb)
dt_predictions_b2.select("prediction", "label", "features").show(25)

# evaluate the model
dt_evaluator_b2 = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(dt_evaluator_b2.evaluate(dt_predictions_b2,
    ↳{dt_evaluator_b2.metricName: "areaUnderROC"})))
```

prediction	label	features
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]

0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]
0.0	1	[2.32168783618879...
0.0	0	[0.00139301270171...
0.0	0	[0.00185735026895...
0.0	0	[0.00336644736247...
0.0	0	[0.00359861614609...
0.0	0	[0.00429512249694...
0.0	0	[0.00464337567237...
0.0	0	[0.00499162884780...
1.0	0	[0.00684897911675...
1.0	1	[0.00835807621027...

+-----+-----+-----+-----+
only showing top 25 rows

Test Area Under ROC: 0.43740745891175686

1.0.8 4.3 Random Forest

```
[41]: from pyspark.ml.classification import RandomForestClassifier

rf = RandomForestClassifier(featuresCol = "features", labelCol = "label")
rfModel_b = rf.fit(trainSetb)
rf_predictions_b = rfModel_b.transform(testSetb)
rf_predictions_b.select("prediction", "label", "features").show(25)
```

	prediction	label	features
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]	
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]	
0.0	1	(24, [0, 1, 2, 3, 4, 5, ...]	
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]	
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]	
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]	
0.0	0	(24, [0, 1, 2, 3, 4, 5, ...]	

```

|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      1|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      0|(24,[0,1,2,3,4,5,...|
|      0.0|      1|[2.32168783618879...|
|      0.0|      0|[0.00139301270171...|
|      0.0|      0|[0.00185735026895...|
|      0.0|      0|[0.00336644736247...|
|      0.0|      0|[0.00359861614609...|
|      0.0|      0|[0.00429512249694...|
|      0.0|      0|[0.00464337567237...|
|      0.0|      0|[0.00499162884780...|
|      0.0|      0|[0.00684897911675...|
|      1.0|      1|[0.00835807621027...|
+-----+-----+-----+
only showing top 25 rows

```

```

[42]: rf_evaluator2 = BinaryClassificationEvaluator(labelCol="label")
accuracy = rf_evaluator2.evaluate(rf_predictions_b)
print("Accuracy = %s" % (accuracy))
print("Test Error = %s" % (1.0 - accuracy))

```

```

Accuracy = 0.7793560340791312
Test Error = 0.22064396592086877

```

```

[43]: rf_evaluator_b = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(rf_evaluator_b.evaluate(rf_predictions_b,
↳{rf_evaluator_b.metricName: "areaUnderROC"})))

```

```

Test Area Under ROC: 0.7793560340791312

```